



University
of Glasgow | School of
Computing Science

Honours Individual Project Dissertation

MQUICTT: EVOLVING IoT AT THE TRANSPORT LAYER

Ivan Nikitin
March 3, 2022

Abstract

The standard protocols used to provide secure communication - TCP and TLS, often fall short for hardware constrained devices due to latency, connection overhead, head of line blocking and large hardware footprints. QUIC - a new transport layer protocol is set to succeed the TCP/TLS stack and is advantageous for performance and security by design. Due to the popularity of MQTT, a lightweight IoT protocol, in this work, we present MQuicTT - a QUIC based implementation of the MQTT protocol. We analyse the landscape of existing QUIC implementations that we considered to underlay MQuicTT, discuss our implementation of a socket API for QUIC and the implementation of MQuicTT itself. We also compare the performance of the resulting solution with existing MQTT implementations in various use cases. Finally, we find that TLS contributes a significant overhead to the binary size of QUIC implementations and discusses possible alternatives.

Education Use Consent

I hereby grant my permission for this project to be stored, distributed and shown to other University of Glasgow students and staff for educational purposes. **Please note that you are under no obligation to sign this declaration, but doing so would help future students.**

Signature: Ivan Nikitin Date: 25 March 2022

Contents

| | | |
|----------|---------------------------------------|-----------|
| 1 | Introduction | 1 |
| 1.1 | Dissertation Outline | 1 |
| 2 | Background | 3 |
| 2.1 | The evolution of the transport layer | 3 |
| 2.2 | The Internet of Things | 5 |
| 2.3 | MQTT | 7 |
| 2.4 | Hardware constrained protocols | 8 |
| 2.5 | Rust programming language | 9 |
| 3 | Requirements and design | 11 |
| 3.1 | Requirements | 11 |
| 3.2 | Design of implementation | 11 |
| 3.2.1 | QUIC | 12 |
| 3.2.2 | MQTT | 14 |
| 3.2.3 | MQuicTT design | 14 |
| 3.3 | Network performance experiment design | 14 |
| 3.4 | Binary size experiment design | 17 |
| 4 | Implementation | 19 |
| 4.1 | QuicSocket | 19 |
| 4.2 | MQuicTT | 20 |
| 5 | Analysis | 23 |
| 5.1 | Connection time comparison | 23 |
| 5.2 | Data transmission comparison | 24 |
| 5.3 | Binary size breakdown | 24 |
| 5.3.1 | MQuicTT binary composition | 24 |
| 5.3.2 | Reducing the QUIC stack | 25 |
| 6 | TLS and possible alternatives | 32 |
| 7 | Conclusion | 33 |
| | Appendices | 34 |
| A | Data | 34 |
| B | MQTT simulation messages | 35 |
| C | Mininet Topology Definition | 36 |
| | Bibliography | 38 |

1 | Introduction

The Internet of Things (IoT) is rapidly becoming the most prominent form of computation, with the number of IoT devices projected to surpass 75 billion by 2025 (Statista 2016). IoT devices range from daily consumer gadgets such as wearables to sensors that are used in smart factories. The central theme of these devices is network connectivity. Network firmware installed on these devices has to be performant enough to send and receive massive amounts of data. Hence, the need for efficient, lightweight network firmware often means that manufacturers of these devices forego security for performance Ling et al. (2018). Therefore, the mass use of devices that are manufactured without security in mind creates a situation where many devices that are critical to infrastructure become vulnerable to cyber attacks, as seen in the 2015 attack on the Ukraine power grid (Liang et al. 2017). However, we can not just employ the usual methods to secure firmware on these devices due to the presented constraints in terms of their physical size, power consumption needs and available hardware. Hence, there is a need to develop methods for secure, performant communication for this device class.

MQTT (OASIS 2014) is a popular message-passing network protocol designed to be lightweight for the IoT use case. While certain protocols have been developed to specifically act as data transfer protocols for IoT devices, MQTT remains the most widely used. MQTT's design ensures that the protocol's implementations have a small code footprint and take up minimal network bandwidth. Significantly, MQTT relies on a transport layer protocol to send data and ensure secure communication. The standard in current MQTT implementations is to use the TCP/TLS stack to provide secure data transfer. As we present in this work, this stack adds overhead, which means that the advantages of MQTT are negated if we need secure communication. However, QUIC (Iyengar and Thomson 2021), a new transport layer network protocol initially designed at Google, is set to succeed TCP with a number of improvements.

QUIC boasts advantages in both secure communication and performance. Expanding the range of hardware constrained devices that MQTT can support would also work towards standardising messaging protocols for IoT devices. As we will demonstrate in further sections, many of the available protocols rely on the same fundamentals, and it may not be desirable to support all of these. Hence, we introduce MQuicTT - a QUIC port of an MQTT library in the Rust programming language, discuss the design choices made during its development, analyse its performance and discuss the challenges IoT presents for network protocols. We further analyse the steps needed to create protocols for hardware constrained devices and present an analysis for the methodology of lowering the overhead of a QUIC stack.

1.1 Dissertation Outline

The rest of this dissertation is structured as follows:

- Chapter 2 provides a background on the recent developments in the transport layer network protocol space, IoT and the Rust programming language.
- Chapter ?? provides a rationale for the implementations of QUIC and MQTT that were chosen for MQuicTT.

- Chapter 4.1 details the implementation of *QuicSocket* – an intermediate QUIC API that we have developed.
- Chapter 3.3 gives an explanation of our methodology for the analysis that follows in later chapters.
- Chapter 5 details the analysis of MQuicTT and provides a discussion of results.
- Chapter 6 discusses the pros and cons of TLS as a security protocol for IoT.
- Finally, Chapter 7 concludes and summarises important results.

2 | Background

In this chapter, we present the background of topics that are fundamental for this work. We examine the current developments in the networks transport layer, the constraints surrounding IoT devices, hardware constrained network protocols, the MQTT protocol and the Rust programming language.

We demonstrate why it is favourable to use QUIC as the transport layer protocol for MQTT and how it can solve some of the challenges in designing protocols for hardware constrained devices. Additionally, we explain why we have chosen Rust as the programming language of choice in the context of creating efficient, secure IoT firmware.

2.1 The evolution of the transport layer

This section presents a background of the network protocols at the transport layer, starting from the currently dominant standards and ending with new advancements in the field. We also consider how to provide secure communication at the transport layer using encryption.

First described by Cerf and Kahn (1974), the Transmission Control Protocol (TCP) has been the primary protocol of the Internet suite since its initial implementation. TCP provides a *reliable* and *ordered* delivery of bytes, ensuring that data is not lost, altered or duplicated and delivered in the same order as intended by the sender. TCP achieves this by assigning a sequence number to each transmitted packet and requiring an *acknowledgement* (commonly referred to as ACK) from the receiving side. If an ACK is not received, the data is re-transmitted. TCP can also use the sequence numbers to order packets in the order intended by the sender on the receiving side.

As TCP is a connection-based protocol, connection establishment must occur before any data can be transmitted. The receiving side (the server) must bind to and listen on a network port, and the sender (the client) must initiate the connection using the process of a *three-way handshake* as shown in Figure 2.1. In the first step of the handshake, the client sends a segment with a *synchronise sequence number* (SYN) that indicates the start of the communication and the sequence number that the segment starts with. The server responds with an acknowledgement - ACK, and the sequence number it will start its segment with - SYN. Hence, we refer to this step as the SYN-ACK. In the third and final step, the client must acknowledge the response. At this point, TCP establishes the connection and can transfer data on it.

In order to achieve *secure communication*, TLS (Rescorla 2018) is often used in the TCP stack. In order to do this, a separate TLS handshake has to occur in order to specify the version of TLS to use, decide on the cypher suites, authenticate the server via its public key and certificate authority's signature, and generate a session key that the protocol uses for symmetric encryption during communication. In the TLS handshake, the first step is for the client to send a *ClientHello* message that specifies the highest version of TLS that the client supports, a list of suggested cypher suites, compression methods and a random number. The server responds with a *ServerHello* message containing the selected TLS version, cypher suite, compression method, and random number. The server then sends its certificate and *ServerKeyExchange* message along with the *ServerHelloDone* message indicating that it has completed its part of the negotiation process. The




Figure 2.1: The TCP handshake needed for connection establishment. The values of the SYN and ACK fields being set indicate the kind of segment sent. For example, for the SYN-ACK stage both the SYN and ACK fields are set.

client will respond with the *ClientKeyExchange* message (CKE), which may contain a public key depending on the chosen cypher suite. Following this, the client sends a *ChangeCipherSpec* message indicating that communication from this point is authenticated and encrypted. The client combines this message with a client finished message. The server responds with the same message, establishing the TLS connection.

Due to the establishment of communication and properties guaranteed by TCP, this process lengthens communication latency. Hence, in use-cases where reliability and connection state is not required, the User Datagram Protocol (UDP) (Postel 1980) is preferred. UDP uses a connectionless communication model with a minimal number of implementation semantics. The only mechanisms provided by UDP are port numbers and checksums in order to ensure data integrity. UDP is preferable for real-time systems as using TCP would cause overhead to latency and retransmission of packets that the application no longer needs.

While TCP and UDP have been the dominant standards of the internet since its creation, QUIC is a relatively new general-purpose transport layer protocol first designed by Roskind (2012) at Google as part of the Chromium web engine and standardised by the IETF in 2021 (Iyengar and Thomson 2021). QUIC aims to improve upon, and eventually make obsolete, TCP by using the concept of multiplexing over a UDP connection. Multiplexing is a method of combining several signals or channels of communication over one shared medium. QUIC makes use of multiplexing by facilitating data exchange on the UDP connection through the concept of *streams*. Streams are an ordered byte-stream abstraction used by the application to send data of any length. Any number of QUIC streams, unidirectional or bidirectional, can be created by either side during the connection. Hence, QUIC allows an arbitrary number of streams to send arbitrary amounts of data on a single UDP connection, subject to the constraints imposed by flow control.

By doing so, QUIC also achieves other performance benefits. For example, QUIC lifts congestion control algorithms from the kernel space to the userspace. Hence, congestion control algorithms

can evolve without being tied down to kernel level semantics and constraints.

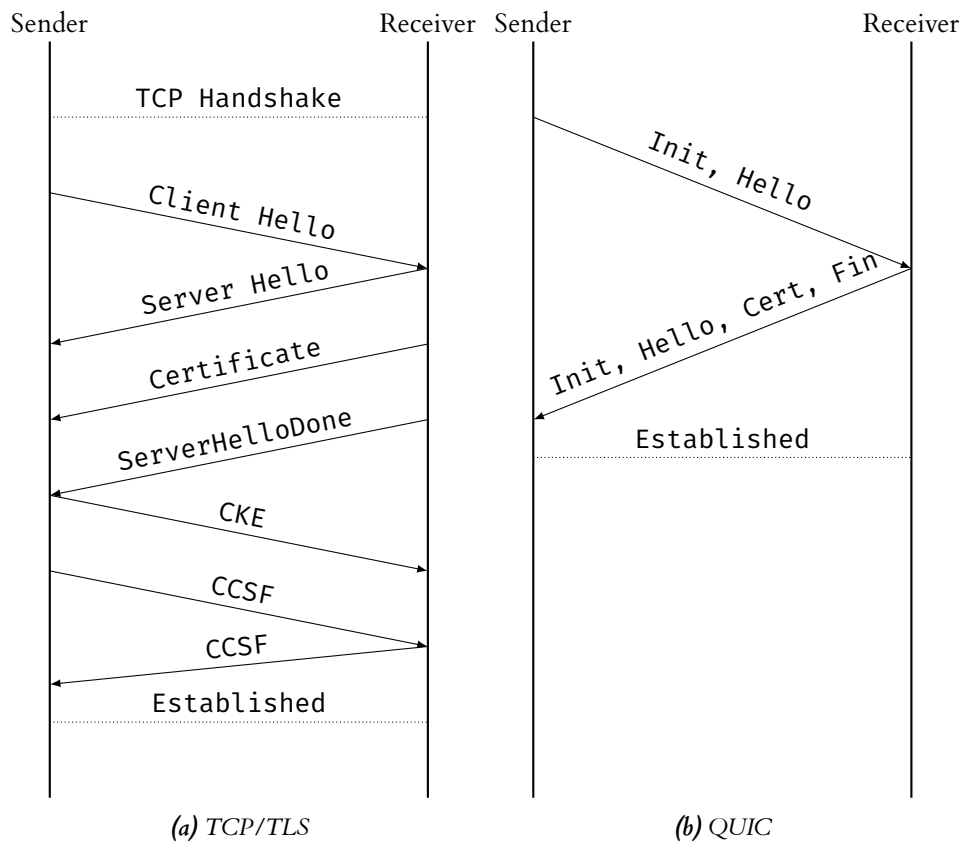


Figure 2.2: Handshakes required to establish secure data transmission in the TCP/TLS stack (a) and the QUIC stack (b). In the case of TCP/TLS, we can see that the handshake is substantially more complex and that the TLS handshake requires the full TCP handshake before it can proceed. In both cases, these handshakes can be made quicker. In the case of TLS, version 1.3 allows for one less round trip before data can be sent, and in the case of QUIC 0-RTT, connection re-establishment may be used in some cases, allowing to send data in the first packet.

Compared to TCP/TLS, QUIC combines the transport and cryptographic handshakes to minimise the time needed for connection establishment; figure 2.2 shows a comparison of the handshakes. QUIC still uses TLS functionality to secure communication as described by Thomson and Turner (2021) unless the developer specifies a different cryptographic protocol; however, this is done differently to TCP. The initial QUIC handshake keeps the same handshake messages as TLS; however, it uses its framing format, replacing the TLS record layer. This use ensures that the connection is always authenticated and encrypted, unlike TLS, where the initial handshake is vulnerable. The combination also means that QUIC typically starts sending data after one round-trip, achieving security by default and lower latency.

2.2 The Internet of Things

It is hard to give a set criterion or definition for which devices qualify as IoT devices. Generally, an IoT device is a small chip that possesses some processing power and may have embedded sensors. The key aspect of IoT devices is that they facilitate data exchange with other devices and systems over the Internet. The modern version of IoT can be attributed to Weiser's (1991) work

on ubiquitous computing, although the term itself first appeared in a speech by Peter T. Lewis in 1985. IoT has applications in various fields, including smart home automation, healthcare, consumer applications, etc.

In terms of classifications within networking and IoT technologies, we can generally split them into wireless and wired, with the former split into short-range, medium-range, and long-range.

Short-range wireless IoT technologies include Bluetooth mesh networks, Z-wave, ZigBee and Wi-Fi, and other lesser-used technologies. These technologies are the most prevalent in consumer IoT devices, such as smart home applications. Medium range networks are used heavily in mobile devices with technologies such as LTE and 5G. The technologies again present an interest due to the amount of traffic that the Internet sees from mobile devices. On the other hand, long-range networks are rather specific in their applications; for example, VSAT – a satellite communication technology that uses small dish antennas is not something we would necessarily think of when encountering IoT. Due to the limited application of long-range technologies, we opted to leave them out of our analysis.

Ethernet remains the dominant general-purpose networking standard in terms of wired technologies used by IoT devices. Although wired technologies provide advantages in terms of data transfer speed, they limit deployments due to the physical wiring constraints.

Having defined what we mean by an IoT device in a network, we now consider the constraints that apply to these devices. Due to the use cases of IoT, the form factor of these devices should be small. For example, a processing unit inside a home assistant has to fit in its enclosure. Additionally, many of these devices have to run for long periods, sometimes on a single lithium battery, hence needing to consume as least energy as possible. Many use cases of IoT devices also require many of them connected in a mesh network. For example, Ericsson (2018) estimated that 0.5 connected devices were used per square meter in a smart factory, with demand growing. Large scale deployments add economic constraints to IoT devices as they need to be manufactured from relatively cheap components.

These constraints mean that IoT devices are limited in hardware resources. Hardware limitations come in three primary forms – CPU power, memory and storage. Storage in the form of flash memory provides the hardest to solve problems regarding secure data transfer. The keys required for protocols such as TLS are often large and need to be stored. For example, the *ESP8266* controller, a widely used IoT chip, comes with 4Mb of flash memory. After installing the firmware and binaries needed, little to no memory may remain for additional storage.

The circumvention of these constraints at the cost of insecure firmware and communication, amongst other issues, is why IoT has become synonymous with security concerns. Efforts to classify the security issues in the IoT space (Alaba et al. 2017; Gupta and Lingareddy 2021; Swamy et al. 2017) and create a taxonomy have generally shown several main topics: insecure firmware level code, issues with privacy due to authentication and authorisation and general security concerns due to poor encryption at the transport layer.

Insecure firmware in IoT devices comes from issues with firmware updates and general vulnerabilities stemming from code. A primary reason is that most programmers opt to create software for IoT devices in memory unsafe languages. Languages such as C provide the needed efficiency to circumvent processing constraints; however also leave room for memory management issues, leading to vulnerabilities. When it comes to privacy, a primary goal is ensuring data integrity and confidentiality. Data sent via the network must not be tampered with nor snooped on by third parties during communication. Man in the middle attacks is a prime concern for protocols such as MQTT due to their lack of self-imposed encryption.

Hence, finding a way to circumvent the hardware constraints presented by IoT devices and still provide secure data transfer is paramount to the safe adoption of IoT. Additionally, opting to

create IoT firmware code in a memory safe language may prevent vulnerabilities that are present on IoT devices from the moment of deployment.

2.3 MQTT

Considering the constraints that apply to IoT devices, we will now look at one of the widely used application-level IoT protocols. MQTT, originally standing for Message Queuing Telemetry Transport, is designed to be a lightweight protocol to transport messages between devices using the publish-subscribe method. MQTT defines two types of participants – the broker and the client. The broker is a server that receives messages from all clients and then routes these messages to the appropriate destinations. On the other hand, a client is just a device that runs an MQTT library and sends the broker messages.

The broker handles the routing of messages in MQTT via *topics*. When a client wishes to publish a message, it does so on a client, and the broker distributes this message to all other clients subscribed to this topic. The broker facilitating communication means that the publishing client does not need to keep track of other clients' locations to communicate with them.

Communication via MQTT can happen after the initial connection request from the client and the subsequent connection acknowledgement from the broker. The connection request can specify a quality of service, referred to as the QoS parameter, to indicate the nature of message delivery.

The possible QoS parameters are as follows:

- At most once (fire and forget) – a client sends a message once, and the broker takes no steps to acknowledge delivery.
- At least once (acknowledged deliver) – a message is re-sent until the broker acknowledges it has received it.
- Exactly once (assured delivery) – the client and broker have to participate in a two-way handshake to ensure that a single copy of a message is received.

We present an example of an MQTT connection with two clients and a broker with QoS of at most once in Figure 2.3. With each level of QoS measure increasing, so does the communication overhead. However, this measure only affects the MQTT part of the communication. The underlying transport layer protocol, such as TCP, will still act as intended. Hence, MQTT relies on an underlying transport-level protocol for data transmission. Additionally, MQTT sends all its connection credentials in plain text format; hence, the communication is vulnerable if the transport layer does not provide encryption. The most widely used suite for providing data transfer and encryption for MQTT is TCP/TLS; however, any protocol that provides lossless, bi-directional communication can be used.

As QUIC provides such a form of communication, we can see that we can use it as the transport level protocol for MQTT. The benefits of this are numerous. For example, perceived performance benefits from lower communication overhead and encryption at header and packet levels. QUIC also comes with the benefit of multiplexing. QUIC can open several streams on a single connection instead of opening several connections from one client. Multiplexing will later play an essential role in our implementation of MQTT/QUIC.

The first implementation of MQTT using QUIC presented by Kumar and Dezfouli (2019) uses the *ngtcp2* implementation of QUIC in the C programming language. The authors find that QUIC reduces the connection overhead time significantly and reduces the processor and memory usage. We present a comparison of these results to the ones obtained through the evaluation that we have conducted in chapter 5.

Due to the authors' chosen underlying QUIC implementation, their API implementation requires its own message queue amongst other common functions such as key settlement. The QUIC




Figure 2.3: The MQTT connection, specifying the QoS measure, followed by the publishing of a message to a broker. It is important to note that the QoS measure does not impact the underlying data transmission provided by a protocol such as TCP. In this example, client B publishes message B to topic A after connecting and the broker publishes this message to client A as it previously subscribed to topic A.

implementation chosen for this project allows the library to handle some of this logic instead. We justify our choice of QUIC library and compare existing implementations in chapter 3.2.1.

2.4 Hardware constrained protocols

We now motivate our focus on MQTT using QUIC by comparing this approach to other popular IoT protocols.

The Constrained Restful Environments (CoRE) working group (CoRE 2016) develops a framework for REST applications that run on constrained devices. The primary protocol that CoRE works on is the Constrained Application Protocol (CoAP) (Shelby et al. 2014). CoAP is a specialised transfer protocol designed to work on constrained devices. CoAP emphasises a generic design that supports low overhead, machine-to-machine interaction, multicast and asynchronous messaging. As CoAP is seen as the successor to HTTP for hardware constrained devices, we do not focus on a comparison to HTTP. Other popular IoT protocols primarily consist of proprietary protocols that are phasing out due to the webs standardisation of REST. Hence, our primary

point of comparison to MQTT with QUIC is CoAP.

CoAP works by providing a usual request/response model between endpoints that mimics the style of HTTP requests while using UDP at the transport layer. In contrast to HTTP, CoAP relies on datagram-oriented UDP communication as its transport to support some of these features. The abstraction that CoAP provides is a UDP layer that deals with asynchrony and a request/response layer that specifies interactions. The header structure for CoAP and MQTT is very similar, with the significant difference being that MQTT needs an ordered reliable packet delivery mechanism. In the past, this restricted MQTT to the TCP/TLS stack; however, with the introduction of QUIC, we can see further similarities between CoAP and MQTT using QUIC with both protocols providing data transfer using UDP. A report on the active deployment of data transport protocols T-Mobile (2018) showed that 62.61% of deployments rely on MQTT as one of their protocols. In contrast, CoAP was the most widely used UDP based protocol, with 22.49% of deployments relying on it as one of their protocols. This is in part because the class of constrained devices that CoAP targets does not cover the entire range of constrained IoT devices. The authors show that MQTT is more suitable for less constrained devices with long-standing connections; however, it may not be suitable for relatively more constrained devices.

In this work, we have focused on providing MQTT with similar advantages that CoAP enjoys due to UDP and analysing if MQTT can be made more efficient to capture the range of more demanding hardware constrained devices. Capturing a wider range of constrained devices using MQTT would increase the number of deployments that can take advantage of its feature set. Additionally, the reasons for our choice stem not only from MQTT's popularity but also due to the exciting prospect of solving its historical dependency on TCP/TLS. This is further reinforced by CoAP showing a similar approach and being widely deployed for hardware constrained devices.

2.5 Rust programming language

Rust is a modern systems programming language created at Mozilla designed to be highly safe and performant. It is a systems language that aims to maintain the performance that we expect from languages like C while also using a unique *ownership* system to maintain memory safety. Instead of garbage collection, Rust opts for a system managed through the resource acquisition is initialisation (RAII) principle (Rust 2021). All values have a unique owner, and their scope is tied to this owner. Hence, by design, Rust does not allow dangling pointers, null pointers, and data races as the compiler will not allow for a programmer to compile unsafe code without circumventing it using the *unsafe* keyword.

For example, consider the program written in C in Listing 2.1 compared to a similar application in Rust in Listing 2.2. The C program demonstrates a use after free bug. The pointer is freed and then used in the print statement, resulting in undefined behaviour. While this example is relatively trivial, these bugs are often tough to debug in a more extensive, more complex system, leading to security vulnerabilities. This issue does not only exist in codebases and organisations with low resources; at the BlueHat security conference, Microsoft researcher Miller (2019) presented that Microsoft targets roughly 70% of their yearly patches at fixing memory safety bugs. On the other hand, the analogous Rust code will not compile due to the ownership system, mitigating this issue altogether.

```
void bar() {
    int *ptr = (Point *) malloc(sizeof(int));
    free (ptr);
    printf("%d", *ptr); // obvious use after free, however this will compile
}
```

```
fn bar() {
    let example = String::from("Example");
    let mut example_ref = &example;
    {
        let new_example = String::from("New Example");
        example_ref = &new_example;
    }
    println!("our string is {}", &example_ref); // causes a compiler error in
    Rust: error 'new_example' does not live long enough
}
```

Listing 2.2: A similar application in Rust will not compile due to the safety guaranteed by the ownership system. A borrow occurs when `example_ref` is assigned to point to `new_example`, however the ownership system recognises that the borrowed value does not live long enough.

Listing 2.1: An example of a use after free in C code. This is an incorrect use of dynamic memory management, however it compiles. In large, complex code bases, missing bugs such as these often happens and can cause exploitable security issues.

When it comes to networked systems in the IoT space, Rust is a natural application due to its focus on concurrency and safe systems programming. Firstly, a memory-safe language may circumvent security-related bugs in IoT firmware and the supporting network stacks. As previously discussed, getting the firmware correct on the first try is essential in IoT due to the difficulty of providing updates. However, assessing the performance of Rust implementations of the QUIC stack is vital to solidifying Rust as a performant systems programming language for hardware constrained devices. If the binary sizes produced by the Rust implementations are larger than their C equivalents or if the code is not as performant, then the memory safety guarantees may not matter for IoT developers.

3 | Requirements and design

The following section presents the identified objectives and the methodology we follow for the rest of the work. We discuss the steps we have taken to analyse the existing QUIC and MQTT implementations to motivate our choice of libraries. Additionally, we describe the test bench created for the network tests of the implementation and the methodology for its binary analysis.

3.1 Requirements

In order to create analyse QUIC as a transport layer protocol for IoT, we have identified the following requirements using a MoSCoW style analysis:

- **Must:**
 - We **must** create a test case for MQTT using QUIC - MQuicTT.
 - We **must** analyse the performance of MQuicTT compared to base MQTT implementations.
 - We **must** focus on realistic testing scenarios for evaluation.
- **Should:**
 - We **should** analyse a general methodology for reducing the hardware footprint of transport layer protocols for IoT.
 - We **should** analyse the hardware footprint and feature set of TLS.
- **Could:**
 - We **could** create a QUIC extension that uses a lightweight security protocol.

From these requirements, we have split the project into the following objectives:

- **O1:** Compare existing QUIC and MQTT implementations that can act as underlying libraries.
- **O2:** Create an abstraction layer for the QUIC specific logic.
- **O3:** Implement MQuicTT by changing the underlying MQTT implementation to use this abstraction layer.
- **O4:** Create a test bench for the network performance analysis of MQuicTT.
- **O5:** Devise a methodology for trimming down the QUIC stack in terms of binary size.
- **O6:** Analyse the components of the MQuicTT binary.
- **O7:** Analyse the features of TLS and possible alternatives.

Hence, the rest of this chapter is split into sections corresponding to the design for each objective.

3.2 Design of implementation

In this section, we present a comparison of existing QUIC and MQTT implementations that can be used as underlying libraries for MQuicTT. Additionally, we discuss the high-level design of MQuicTT and the approach we have taken when it comes to using existing libraries compared to creating new protocol implementations.

3.2.1 QUIC

In this section, we compare the existing implementations of the QUIC protocol in the context of usability for IoT devices and present the reasoning behind our choice of QUIC library. We have considered the mainstream implementations in the C, C++, Rust and Go programming languages as these are the languages that can be considered systems languages and would thus be the most widely used ones for IoT devices. In addition to this, we did not consider implementations paired with a browser web engine as these would be impossible to use on hardware constrained devices. Hence, we did not include notable implementations such as the QUIC implementation of the chromium web engine (Chromium 2021) and *Neqo* (Mozilla 2022).

Importantly, to analyse the performance of underlying QUIC implementations, we have transferred the same file using each implementation several times and have taken the average time to do so. We repeated this while altering parameters as described in Section 3.3. The differences in transfer time were negligible and have hence not been presented as part of the comparison of libraries. Hence, we compare QUIC implementations in terms of their TLS use, programming language and the binary size produced by their client and server implementations.

First, we consider the general landscape of available QUIC implementations to contextualise the ones developed in Rust. Table 3.1 demonstrates the analysed QUIC implementations. In order to find the size of the binary, we have used the Linux *ls* utility. In the case of the *mvfst* implementation, we have taken further steps due to the reported binary size being far too large. Therefore, we had to remove the C++ debug symbols that were causing the binary size to be over 300 MiB. We have identified the following main methods by which QUIC implementations incorporate TLS:

- Use of an external library – the implementation uses an external TLS API either by using a package manager in the case of Rust or by relying on an installed implementation in the case of C and C++.
- Use of own implementation – the implementation packages its implementation of the TLS protocol alongside QUIC.

This different use of TLS presented a challenge in calculating the binary size of the QUIC servers and clients.

Table 3.1: The identified QUIC implementations and their binary size footprint. The footprint has been split into a client and server footprint using a minimal reproducible example for each. We used each implementation to create an example client and server capable of sending and receiving QUIC packets and analysed the binary size. Where provided, we compared this to the given examples to ensure as little implementation bias as possible. However, this is still not a perfect estimate, and some variance due to implementation details may be present.

| Implementation | PL | Footprint (client) (MiB) | Footprint (server) (MiB) | TLS method |
|----------------|------|--------------------------|--------------------------|------------|
| ngtcp2 | C | 3.4 | 4.3 | External |
| picoquic | C | 3.2 | 3.9 | External |
| msquic | C | 3.2 | 4.1 | External |
| quic-go | Go | 8.7 | 9.9 | External |
| Quinn | Rust | 9.1 | 9.5 | External |
| Quiche | Rust | 7.8 | 7.0 | Own |
| mvfst | C++ | 11.1 | 12.0 | Own |

Specifically, in each case, we have considered the external dependencies that a developer will have to install to run the QUIC implementation on a device. Hence, in the case of C implementations that require an external TLS library, such as OpenSSL, to be installed and linked on the system, we have opted to add the binary size produced by OpenSSL to the size of the QUIC binaries.

Additionally, in the case of *picoquic*, we have added the size of the *picotls* dependency on top of the size of OpenSSL.




Figure 3.1: The sizes of the client and server footprints for the selected QUIC implementations. Notable, only Quiche produced a server binary with a smaller size than its corresponding client binary. It is hard to estimate the error margin for the data as this depends on implementation details in the example client and servers.

Figure 3.1 further visualises the comparison of binary sizes between the various implementations. This is an important aspect due to the aforementioned hardware constraints. Notably, we can see that five out of the seven analysed implementations opt to use an external TLS library or engine. Out of these five, all C implementations supported OpenSSL, with the Go and Rust implementations opting to use a TLS library. In the case of *quic-go*, this was the *crypto/tls* package, and in the case of Rust, *rustls*. We can also see that the Rust implementations are not drastically different in binary footprint size.

In terms of using a memory-safe language, although Go is described as memory safe, it does not opt for compile-time memory safety and instead uses the panic model. The panic model was the primary reason for the choice of Rust instead of Go, as Rust provides these guarantees using a robust type system. Between the two Rust implementations – Quinn and Quiche, we chose Quinn due to issues with the Quiche library. On the other hand, Quiche opts to make the user create a *mio* event loop, which interfered with the *tokio* runtime environment used in our chosen MQTT implementation discussed in the next section. In addition to this, we found that the Quinn API is easier to work with when creating the intermediate library discussed in Chapter 4.1; for example, Quinn handles the QUIC handshake in the library and does not require the developer to create an event loop.

3.2.2 MQTT

Compared to QUIC, the choice of MQTT implementation was substantially simpler. The criteria for MQTT implementation were that it was developed in Rust, implemented both a client and a broker, had widespread use and adhered to MQTT version 5.0. Based on the above criteria we identified two possible implementations: Eclipse’s (2018) *paho* and *rumqtt* (Bytebeam 2020). We opted to use *rumqtt* as it is a native Rust implementation, whereas *paho* provides a rust binding to an underlying C implementation. We chose to evaluate a fully Rust native MQTT/QUIC stack as this provides an opportunity for a valuable comparison to mainstream C implementations. Other available implementations supported only one side of the MQTT protocol or only supported version 3.1.1.

Rumqtt provides components for an MQTT application – *rumqttd* and *rumqttd*. The former can be used to create an MQTT client, and the latter a broker. However, the code base for these is somewhat similar, easing the incorporation of QUIC. Both components provide an interface for supporting asynchronous communication using a *tokio* runtime, which fits nicely into our choice of QUIC implementation as *Quinn* requires a *tokio* environment. By default, *rumqtt* uses TCP as its transport layer protocol and TLS through *rustls*. All underlying implementation assumptions remain equal because this is the same library as *Quinn* uses.

3.2.3 MQuicTT design

To create MQuicTT, we have, as previously seen, opted to use existing implementations of QUIC and MQTT. An alternate approach to this would be creating implementations from scratch that could be optimised for our use case. However, we demonstrate an actual use case by opting to use existing libraries. These libraries have both been built for production use and are relatively widely deployed. Hence, we are more likely to get accurate results. In addition to this, the development time needed to create implementations for these protocols from scratch would far exceed this project’s scope.

However, we have created an intermediate library that will act as an abstraction from the QUIC logic – *QuicSocket*. By creating this intermediate library and ensuring that its API mimics a standard TCP socket, we lower the difficulty of porting *rumqtt* to QUIC. Additionally, this library can be used for any application that uses the QUIC protocol.

Hence, the design for the implementation has been split into two broad stages. In the first stage, we create an abstraction around *quinn* that acts as a socket for the QUIC protocol. In the second stage we use the resulting socket library to port *rumqtt* to QUIC.

Further discussion of specific technical choices that we have made during the implementation can be found in Chapter 4.

3.3 Network performance experiment design

We now discuss the methodology for the performance portion of the implementation evaluation. That is, this section focuses on evaluating MQuicTT against the base *rumqtt* implementation.

The critical consideration for this design is the scenarios in which IoT devices are used. When evaluating the network performance of the implementations, we considered two options: using real IoT devices or using a network simulation tool. Due to technical limitations that came with using real devices, such as not being able to access the router of our network, we opted for simulation. In this section, we will discuss how we used Mininet (Lantz and Heller 2013), a realistic virtual network, in our evaluation.

Mininet is a tool that network developers and researchers can use to create software-defined networks (SDNs) using the *OpenFlow* standard.




Figure 3.2: The resulting Dumbbell topology with the broker being $h9$ – the central node. The link between the central switch and the central host results in a congested network.

Using the Python API provided, we created the network topology shown in Figure 3.2. The script takes several parameters to create the following three scenarios:

- A synthetic scenario testing the limits of implementations.
- A realistic scenario based on a smart-home use-case.
- A realistic scenario based on a 3D printer farm.

The topology for all three scenarios remains the same – a minimum spanning tree of a typical IoT mesh network. When designing this topology, we needed it to reflect various realistic IoT scenarios. We can imagine this as a singular room in a home with all of its smart appliances being a cluster in the smart home or a cluster of 3D printers in a section of a farm in the 3D printer farm. The critical feature of this topology is that many clusters are connected to a central broker that controls the topology, hence creating a dumbbell structure. The full definition of this topology can be found in Appendix C. The variables that the script changes between scenarios and simulations are the link's *bandwidth*, *delay* and the rate of *packet loss*.

The bandwidth of a link is the maximum rate of data transfer we can achieve. In contrast to bandwidth in signal processing, we measure bandwidth in bits per second rather than hertz in computer networking. The delay of a link specifies the latency of the link. It is the time that a bit of data takes to travel across a link. We measure this in milliseconds. Link delay corresponds to the geographical distance between the communicating parties; however, in the case of IoT, we can expect devices to be in local proximity. Lastly, the packet loss rate shows the percentage of corrupted or dropped packets in transit. Various protocols having to retransmit packets also adds to the delay of data transfer. Importantly, we have only considered the typical circumstances of packet loss and have not included scenarios such as interference or packet loss attacks.

The bandwidth and delay numbers correspond, as closely as possible, to various link types in a network. To do so, for the smart-home scenario, we have gathered data from the Ofcom (2021) report on UK broadband speeds. There were specific cases in which it was not possible to find this data in the report; hence it was augmented using a similar methodology in work conducted by Prevdi et al. (2019) and in the case of ZigBee, the work by Alena et al. (2011).

Table 3.2: The parameters chosen for each link simulation in Mininet in the smart home scenario. The types of links were chosen as the most commonly occurring ones in IoT use cases.

| Simulated Link Type | Link bandwidth (Mb/s) | Link delay (ms) | Packet loss rate (%) |
|---------------------|-----------------------|-----------------|----------------------|
| Wi-Fi | 30 | 10 | 2 |
| ZigBee | 0.25 | 5 | 1 |
| 4G | 4 | 20 | 1.5 |
| 3G | 1 | 40 | 1.5 |
| 100Mb Ethernet | 100 | 1 | 0.2 |

Hence, we expect that the bandwidth and link delay numbers accurately represent a real-world scenario. However, it was complicated to find exact estimates for packet loss rates, with most sources describing approximations for a stable connection (SDU 2013) and not precise measurements. Hence, the data are best estimates, cross-validated through the different sources and are not exact values.

In the case of the smart home scenario, as presented in Table 3.2, we expect that our packet loss rates are accurate as these, as previously stated, do appear in the home broadband reports and other studies. In the case of the 3D printer, as presented in Table 3.3 farm scenario, due to the number of machines and the interference these cause in a 3D printer farm, we modelled this scenario to have more packet loss. We specifically chose this scenario to test the benefits that QUIC should receive in an environment with high packet loss. However, we found it difficult to find exact data on packet loss percentages in smart factories or workshops; hence, we expect a margin of error on some of the values.

In general, we can assume that the packet loss rates will increase by some constant factor across all links. Hence, we have also created a synthetic scenario with extreme packet loss as presented in Table 3.4. If the protocol performs well for an extreme scenario, we can expect it to also perform well for packet loss percentages up to that scenario.

Table 3.3: The parameters chosen for each link simulation in Mininet in the 3D printer farm scenario. The data assumes a typical IoT setup where most devices are within local geographical proximity. That is, the devices are communicating with each other within the range of one factory or site, with only the central node communicating with some server.

| Simulated Link Type | Link bandwidth (Mb/s) | Link delay (ms) | Packet loss rate (%) |
|---------------------|-----------------------|-----------------|----------------------|
| Wi-Fi | 30 | 10 | 5 |
| ZigBee | 0.25 | 5 | 3 |
| 4G | 4 | 20 | 2.5 |
| 3G | 1 | 40 | 2.5 |
| 100Mb Ethernet | 100 | 1 | 0.5 |

We evaluate MQuicTT's performance using the presented topologies and respective simulation parameters in the following way. Each protocol transmits 100 messages from a client to a broker for each data link.

During this, we measure the following:

- The time taken for the underlying transport protocol to establish a connection.
- The connection time. That is, the time before MQTT can send its first data packet.
- The time taken to transmit the message to the broker.

This is repeated 10 times, and the average transmission time is taken to ensure accurate results. This process is repeated for all three testing scenarios.

Table 3.4: The parameters chosen for each link simulation in Mininet in the synthetic scenario. Here we opt for extreme packet loss scenarios to test the boundaries of the protocols. We opted for 20 times the normal packet loss that the link can expect in each case. We have also tested this for higher values however all protocol implementations greatly suffered in performance beyond this.

| Simulated Link Type | Link bandwidth (Mb/s) | Link delay (ms) | Packet loss rate (%) |
|---------------------|-----------------------|-----------------|----------------------|
| Wi-Fi | 30 | 10 | 15 |
| ZigBee | 0.25 | 5 | 15 |
| 4G | 4 | 20 | 20 |
| 3G | 1 | 40 | 20 |
| 100Mb Ethernet | 100 | 1 | 4 |

In general, MQTT allows for messages with a maximum size of approximately 260MB. However, this is a huge message, and most publicly deployed brokers reject it, so we created a representative message for each use case. Each topic in MQTT consists of a hierarchy of topic levels separated by a forward slash. For example, in a smart home scenario, we may have a topic like *home/groundfloor/kitchen/temp* to control the temperature in the kitchen via a smart thermostat. A topic may also include a wildcard. The topic string *home/groundfloor+/temp* includes a *single-level* wildcard that will match an arbitrary string. This would match the topic *home/groundfloor/lounge/temp*, but not match the topic *home/secondfloor/kitchen/temp*. If a client wishes to subscribe to multiple topics with the same prefix, a *multi-level* wildcard may be used. For example, the topic *home/secondfloor/kitchen/#* can be used to subscribe to all topics with a prefix matching the string before the hash character. Notably, brokers reserve topics for system messages starting with the \$ character.

The chosen topic and the transmitted message for each scenario can be found in Appendix B.

3.4 Binary size experiment design

The last section of this chapter describes how we have conducted further analysis on the binary size of the QUIC protocol that underlines MQuicTT. The focus of this section is to see if the QUIC stack contributes a significant overhead to MQuicTT's binary size and how we can reduce this overhead. We expect that the QUIC stack will contribute to the majority of the size of MQuicTT as MQTT itself is designed to have a low code size overhead. Hence, the first step of this part of the analysis will be to determine how much QUIC contributes to the binary size.

In order to get a breakdown of the binary we have used the *cargo - bloat*¹ utility. The utility analyses the binary using custom ELF, DWARF and Mach-O parsers and disassembles the binary to look for references and links to anonymous data. Doing so creates a map of the binary that shows where every byte has a label attached to it.

This utility provides the composition of a Rust binary. However, it is not perfect and results in some margin of error. Unfortunately, this margin of error is also not easily measurable. By comparing the total size of the binary as reported by *cargo - bloat* to the size reported by the operating system, we have deduced that the total error margin is within 1% with good precision. This should mean that we can get a somewhat accurate error margin on the components. However, it is also possible that the internal calculations are inaccurate despite the overall size being accurate.

The next step in this stage will be analysing methods for trimming down the QUIC stack binary size. Hence, in this stage, we shift our focus to the binary produced by *QuicSocket*. To reduce the

¹cargo-bloat - <https://lib.rs/crates/cargo-bloat>

binary size, we opt to use the method established by Eggert (2020) as recreating these steps may show a general framework for reducing binary sizes for hardware constrained devices. Notably, our application already handles client and server code separately; the MQTT broker requires a different binary to the MQTT client.

Hence, the steps we take are as follows:

- Compile the binary for a 32-bit target by setting the *target* flag in cargo to *i686-unknown-linux-gnu*.
- Remove any error handling code beyond what is needed for the binary to compile.
- Remove any code that writes to standard output.

After every step, we record the difference in binary size made by the change using the same methodology.

4 | Implementation

We now present the implementation for the intermediate QUIC library that we created and MQuicTT. We discuss the specific implementation choices that we have made during development and present example uses of each implementation.

4.1 QuicSocket

This section will present *QuicSocket*¹ – the intermediate wrapper API that we developed using *Quinn* to facilitate establishing a QUIC connection, control streams, and send and receive data.

QuicSocket provides two main types to communicate with *Quinn* – a *QuicServer* and a *QuicClient*, and a trait that both of these implement – *QuicSocket*. The *QuicSocket* trait defines three functions: *new*, *send* and *receive*. When initialised using the *new* function, the server and the client attempt to open a QUIC connection. In the case of the server, it binds to a given port and starts listening for a connection request. The client binds to a random available port and attempts to connect to a provided URL. Additionally, in the *new* function, the server generates a new TLS certificate or reads it from a specified path. The client's *new* function must be supplied with the corresponding certificate's path. If either certificate is not present, the server will reject the connection attempt. Once a connection has been opened, data transfer can happen.

```
async fn send(&mut self, payload: Vec<u8>) -> Result<()> {
    let (mut send, _) = self
        .connection
        .open_bi()
        .await
        .map_err(|e| anyhow!("failed to open stream: {}", e))
        .unwrap();
    send.write_all(&payload)
        .await
        .map_err(|e| anyhow!("failed to send request: {}", e))?;
    send.finish()
        .await
        .map_err(|e| anyhow!("failed to shutdown stream: {}", e))?;
    Ok(())
}
```

Listing 4.1: The *send* function that the *QuicClient* and *QuicServer* use for sending data. Data transfer is facilitated by opening a bidirectional stream on the pre-existing connection.

In order to send data, either side can use the *send* function demonstrated in Listing 4.1. We first open a bidirectional stream and then write the provided payload. Once the client completes writing the payload, it shuts down the stream. On the other side, the receive function in Listing 4.2

¹QuicSocket – <https://github.com/Apolexian/QuicSocket>

```

async fn recv(&mut self, buf: &mut [u8]) -> Result<usize> {
    let (_, mut recv) = self.bi_streams.next().await.unwrap().unwrap();
    let len = recv
        .read(buf)
        .await
        .map_err(|e| anyhow!("failed to read response: {}", e))?;
    Ok(len.unwrap())
}

```

Listing 4.2: A similar application in Rust will not compile due to the safety guaranteed by the ownership system. A borrow occurs when `example_ref` is assigned to point to `new_example`, however the ownership system recognises that the borrowed value does not live long enough.

processes the next available QUIC stream and reads the data on it into a supplied buffer. It then returns the number of bytes written to the buffer. Hence, to use *QuicSocket*, we must either provide or generate TLS certificates, open a QUIC connection and then use the send and receive functions. The developed API is analogous to how a TCP socket sends and receives data. An alternative approach that we considered was to make the *recv* function return the value that the client or server read from the stream. However, this approach does not integrate well into any protocol implementation due to the general preference for the buffer pattern.

Error handling is handled via the *anyhow* package throughout the library for idiomatic error handling. In addition to this, we wrap return values in Rust’s *Result* construct. Hence, all errors are idiomatically mapped wherever a failure can occur.

An important design choice that we have made when developing *QuicSocket* is using QUIC streams. Every call to send and receive in the current implementation creates a bidirectional stream. An alternate choice was to open a stream at the start of the connection and use it until the connection closes. However, the latter approach runs into issues when it comes to reading from the stream. All data was read on one call to receive, which unnaturally coalesced MQTT packets.

Notably, all operations happen asynchronously. However, as of the implementation time, the Rust standard library does not support asynchronous traits. The workaround to this issue is using the *async – trait* crate. An unfortunate drawback of using this crate is that every function call results in a heap allocation due to the crate’s implementation semantics. The extra heap allocations do not usually present an issue; however, in the case of hardware constrained devices, this may lead to a bottleneck depending on the usage of the library. The reasons for not supporting asynchronous traits in Rust are somewhat complex. For one, an asynchronous function in Rust returns an *implFuture*, meaning that an asynchronous trait would have to support returning *implTrait*, which it does not. However, the *async – trait* crate instead returns a *dynFuture*, the Rust implementation of dynamic dispatch, resulting in a heap allocation but allowing it to be used inside of traits. We will not further discuss the other reasons for this crate’s existence; however, Matsakis (2019) created a comprehensive analysis of this issue. Once Rust incorporates asynchronous traits into the language, the heap allocation issue will be resolved, resulting in this problem no longer existing for hardware constrained devices.

4.2 MQuicTT

This section describes how we used *QuicSocket* to create a QUIC port of *rumqtt*, the design decisions that we have taken, and possible alternate approaches.

As previously mentioned, *rumqtt* provides to APIs: *rumqttd* for creating MQTT clients and *rumqtd* for creating brokers. Hence, each of these APIs relies on an underlying TCP implementation to transmit data. A configuration file is also required for *rumqtd* that specifies various parameters such as paths to TLS keys and the ports that MQTT will operate on. Additionally, *rumqtt* implements its own layer of TLS to ensure secure communication on top of the TLS layer provided by the underlying transport protocol.

To create *MQuicTT*, we have taken steps that can broadly be summarised in two phases: altering the default TLS code and changing the underlying transport protocol from TCP to QUIC.

As we are working in the domain of hardware constrained devices, providing multiple layers of encryption would be excessive, even for the goal of secure communication. We can afford for only the transport layer to handle encryption. Hence, the first step that we have taken is to remove all TLS functionality from both *rumqtd* and *rumqttd*.

After this, we identified that both APIs have a central interface that controls network communication. In the case of *rumqttd* this is the *network* struct inside *framed.rs*² and in the case of *rumqtd* it is the *network* struct inside *network.rs*³. Both of these interfaces use a *tcpstream* to send and receive data. The *tcpstream* is opened when the struct is initialised and closed when the MQTT connection ends. As TCP is a stateful protocol, minimal connection handling occurs after opening the initial stream.

```
#[tokio::main(worker_threads = 1)]
async fn main() {
    ...
    let quic_client = QuicClient::new(None).await;
    let mut mqttoptions = MqttOptions::new("test-1", ip, 1883, addr, quic_client);
    mqttoptions.set_connection_timeout(10);
    mqttoptions.set_keep_alive(5);

    let (client, mut eventloop) = AsyncClient::new(mqttoptions, 100);
    task::spawn(async move {
        requests(client).await;
    });
}
```

Listing 4.3: An example of initialising an *MQuicTT* client. The QUIC connection is established by initialising a *QuicClient* and the resulting client is passed as an MQTT option.

As we have modelled *QuicSocket* to have an API that closely resembles a typical *tcpstream* API, the challenge of this stage of the port was managing QUIC streams and the state of connection. Due to UDP, the protocol underlying QUIC, being stateless, we have handled the connection differently in brokers and clients.

The connection request is sent before an MQTT client object is initialised in the client's case. Upon successful connection, the programmer must pass the connection to the initialisation function of the client. This means that a single QUIC connection underpins the client's entire MQTT connection, and streams are used for packets. Hence, when the client wishes to close the MQTT connection, it also closes the underlying QUIC connection. In order to do this we have modified the parameters needed to create an *rumqtt* MQTT client. An example of the initialisation of a QUIC connection and MQTT client using our implementation is demonstrated in Listing 4.3.

²framed.rs - <https://github.com/bytebeamio/rumqtt/blob/master/rumqttd/src/framed.rs>

³network.rs - <https://github.com/bytebeamio/rumqtt/blob/master/rumqtd/src/network.rs>

In the case of a broker, we must wait for a client to send a connection request. Hence, we may initialise the broker and wait for a QUIC connection, only allowing data transfer when any connection is established. This also means that the broker does not need to track which connection comes from which client.

```
fn main() {
    pretty_env_logger::init();
    let config: Config = confy::load_path("rumqtttd.conf").unwrap();
    let mut broker = Broker::new(config);

    let mut tx = broker.link("localclient").unwrap();
    thread::spawn(move || {
        broker.start().unwrap();
    });
    ...
}
```

Listing 4.4: An example of initialising an MQuicTT broker. We can see that no operations with QuicSocket are required for this initialisation as all QUIC operations are handled internally.

Hence, as we can see in Listing 4.4, the programmer is not required to use any *QuicSocket* code when creating an *MQuicTT* broker.

We considered an alternate approach to refactoring the *rumqtt* codebase to consolidate all network code into a shared internal package. This would mean that the API for creating a client and broker would be identical and perhaps more ergonomic. However, we decided not to go with this approach because *rumqttc* and *rumqtttd* would no longer be disjoint, independent components. With the chosen approach, it is possible to use different underlying transport protocols for either side, which allows for flexibility in implementation.

Another avenue of discussion stems from *rumqtt* having two different client implementations: an asynchronous client and an asynchronous client. In either case, a *tokio* eventloop is used to handle events. As *QuicSocket* requires an asynchronous environment, the natural choice was to use the asynchronous client. It could however be possible to handle asynchronous events more efficiently by tying the QUIC eventloop to the MQTT eventloop as in the approach described by Kumar and Dezfouli (2019). In our case, we leave the comparison between these two approaches as an avenue for future work.

5 | Analysis

This section presents the analysis results conducted using the previously described methodology. Each section presents the results, and a discussion follows.

5.1 Connection time comparison

In this section of the analysis, we focus on connection time. By connection time, we mean the time it takes for the underlying transport protocol to establish a secure connection. For this analysis, we have not considered connection re-establishment.

Connection time is essential in IoT devices as many do not keep a connection alive while idling. A device may opt not to be connected at all times to reduce energy consumption and computational power. The device will then establish a connection and transmit data whenever required. However, this means that some efficiency is lost if the connection has to be constantly re-established. Hence, MQTT defines the period to keep the connection alive as the *keep alive interval*. In an MQTT/TCP implementation, the standard way to extend the keep alive interval is to periodically send a *ping* packet, forcing the connection to remain open. However, it has been shown that this method can lead to security vulnerabilities (Vaccari et al. 2020; Mileva et al. 2021).

We have talked previously about QUIC having a less complex handshake and hence being able to establish a secure connection quicker than TCP/TLS. Based on this, we hypothesise that:

- **H1:** *MQuicTT* has a lower connection time across all scenarios.
- **H2:** *MQuicTT* is more resilient to packet loss i.e. as packet loss increases, *MQuicTT* is affected less than the base *rumqtt*.

To measure the connection time, we have followed the previously described methodology. That is, we began a communication via *MQuicTT* and the other MQTT implementations and captured it using *tcpdump*. We have then extracted the time between the start of the QUIC or TCP communication and the handshake completion. The results of this are shown in Figure 5.1.

Based on these results, we can see that *MQuicTT* does not present any advantage in the IoT home scenario. Its connection time is on-par in the cases of ethernet and Zigbee, and considerably slower than the base *rumqtt* in the case of 3g, 4g and WiFi. Although *MQuicTT* performs better than *rumqtt* in other scenarios, these results do not support hypothesis **H1** and indicate that *MQuicTT* does not connect faster in all scenarios. This could perhaps be attributed to QUIC potentially performing worse in low bandwidth environments as discussed by Nepomuceno et al. (2018). The authors attribute QUIC's worse performance to the current infrastructure being optimised to work with existing transport protocols and the relative immaturity of QUIC. However, it is unclear if this would affect results from our test bench based on an SDN.

On the other hand, we can see that the results support **H2** as *MQuicTT* connects faster than the base implementation in scenarios with high packet loss. In fact, the base implementation did not manage to establish a connection in the synthetic scenario for any data link except ethernet. By not being able to establish a connection, we mean that the timeout value of the protocol would have to be increased, so the communication can no longer be considered performant in this case.

In the case of the printer farm scenario, *MQuicTT* outperformed the base implementation across all data link layers. We can also see that the most significant benefit of using *MQuicTT* comes from using it in less reliable scenarios. This follows logically from QUIC's better handling of packet loss.

In summary, in scenarios with requirements for performance, high packet loss and network congestion, *MQuicTT* presents better results in terms of connection time and provides better resilience to extreme packet loss values. However, it may be better to use current implementations for networks that do not experience high packet loss and congestion.

5.2 Data transmission comparison

We next evaluate the time it took both implementations to transmit the aforementioned MQTT messages. To do so, we have used the previously described methodology.

Based on the analysis of connection time and the previous discussion, we hypothesise that:

- **H1:** *MQuicTT* performs on-par with *rumqtt* in the IoT home scenario but does not present a significant difference.
- **H2:** *MQuicTT* performs significantly better than *rumqtt* in the printer farm scenario.
- **H3:** *MQuicTT* can transmit all the messages on all data links in the synthetic scenario, whereas *rumqtt* is not.

Figure 5.2 shows the results of this experiment.

We can see that the results support hypothesis **H1**. That is, *MQuicTT* performs almost identically to *rumqtt* across all data links for the IoT home scenario. In fact, *MQuicTT* actually performs marginally better than *rumqtt* in this scenario with the transmission time of *MQuicTT* being on average 2.02% lower than that of *rumqtt*. From this we can see that *MQuicTT* presents no disadvantages in terms of transmission time even in network with low packet loss and congestion.

Hypothesis **H2** however is not supported. The performance advantage in terms of transmission time for *MQuicTT* in the printer farm scenario is almost identical to that of the smart home scenario. On average, *MQuicTT* transmits the data 2.36% faster than *rumqtt* and although this is an improvement, it is comparable to the improvement in the IoT home scenario.

The result means that the presented use-case does not present a high enough packet loss percentage for *MQuicTT* to have a significant advantage. This can be verified by looking at the synthetic scenario, which validates hypothesis **H3**. In the synthetic scenario only *MQuicTT* managed to transmit the data on all data links and transmitted the data 13.8% faster using ethernet. Hence, in environments that may experience extreme data loss, *MQuicTT* presents a clear advantage, however, finding such a scenario may be difficult.

Overall, *MQuicTT* transmitted the data faster than *rumqtt* across all scenarios and has shown to be more resilient against high packet loss. It may however be ineffective overhead to migrate existing deployments to it due to the marginal benefits in low congestion environments.

5.3 Binary size breakdown

5.3.1 MQuicTT binary composition

We first look at an in-depth breakdown of the *MQuicTT* binary and analyse the contribution to the QUIC stack. Figure 5.3 shows the result of the binary breakdown by crate.

When analysing the QUIC stack contribution to the binary size we can see that *quinn* contributes to 17.8% of the binary size and *rustls* contributes a further 12.6%. Additionally, we can see that

the *ring* crate contributes another 5.8%. This crate is a Rust binding for BoringSSL primitives; hence it is part of the TLS implementation. Overall, this means that the QUIC stack constitutes 36.2% of MQuicTT's client binary size. This result is in line with our prediction that the QUIC stack will contribute the most to the size of the binary.

Unexpectedly, however, the *regex* crate contributes 14.3% to the size of the client binary exceeding even the contribution made by the TLS implementation. This can partially be attributed to Rust's support for Unicode in strings. Specifically, a *char* type in Rust represents a Unicode scalar value, a Unicode version agnostic type. While providing many wanted features, Rust's native support for Unicode also increases the used resources. A possible workaround would be to add optional support for the *unicode :: Ascii* type to the *regex* crate. Another possible reason for the size of this crate is the source codes extensive use of the *inline(always)* annotation. This annotation tells the compiler to inline functions across crate boundaries, possibly contributing to the binary size. Overall, we can see that this anomaly can be attributed to the choice of programming language, which presents an insight into Rust as a language for hardware constrained devices.

We can also see that despite other crates not contributing a considerable amount towards the binary size, cumulatively, their size adds up to 17.3% of the binary size due to the large number of them. We have found over 40 crates that contributed to this, including data structure implementations, lower-level network interfaces, abstractions on byte buffers and logging utilities. Notably, we have found that logging utilities and error reporting do not largely contribute to the size of the client's binary. This comes in contrast to one of the steps in reducing the binary size taken by Eggert (2020).

Figure 5.4 shows similar results for MQuicTT's broker. We can see that in the case of the broker, the QUIC stack constitutes 32.9% of the binary size. Similarly to the client, the *regex* crate again shows a large binary size footprint, and the lesser crates contribute to a sizeable cumulative amount. Additionally, similarly to the results presented for the client, we have found that the logging and error checking crates have not contributed largely to the binary size.

We have further analysed the binary to see which methods contribute to the size for the client and broker. Notably, in both cases, the *quinn_proto* module of *quinn* contains several methods such as *connection :: Connection :: process_payload* and *endpoint :: Endpoint :: handle* that each contribute around 0.8%. Notably, however, the *regex :: exec :: ExecBuilder :: build* method from the *regex* crate contributes 0.7%, which is a larger contribution than *quinn*'s methods for decrypting packets and polling for connections.

5.3.2 Reducing the QUIC stack

As we can see from the composition of the binaries of MQuicTT, the QUIC stack is around a third of the size of the implementation. Hence, we now analyse methods for reducing QUIC's contribution following the previously described methodology.

Figure 5.5 demonstrates the reduction in binary size at each stage described in the methodology. Notably, compiling the binary to a 32-bit only target contributes significantly to the reduction of the binary size. The 32-bit only version of the broker is 5.8% smaller and the 32-bit only client is 8.5% smaller. However, the two remaining stages did not contribute significantly to the reduction of the binary sizes. This can perhaps be attributed to implementation details as the reduction due to error handling and standard output use would be directly proportional to their use in the code base. This is further reinforced by the client being reduced by a higher percentage from removing error handling due to the client's source code having more of this feature.

Hence, from the binary analysis conducted we can see that the QUIC stack contributes a large amount to the overall binary size of MQuicTT. We can also see that reducing the size of the

QUIC stack is not trivial with best results coming from compiling to a 32-bit only target.

When it comes to reducing the size of the TLS component Eggert (2020) attempts create a minimal cypher implementations used by TLS, and we instead opt to discuss the possibility of a complete alternative to TLS, suitable for hardware constrained devices in Section 6.




Figure 5.1: The time it took for the underlying transport protocol to establish a secure connection. The base rumqtt implementation is shown in green and MQuicTT in orange. Each column represents a scenario and each row a data link. The time for each figure is shown in milliseconds. The base implementation of rumqtt could not establish a connection in the synthetic packet loss scenario.

Time taken to transmit all data in base rumqtt vs MQuicTT.



Figure 5.2: The time it took for the implementations to transmit the MQTT messages. The base rumqtt implementation is shown in green and MQuicTT in orange. Each column represents a scenario, and each row is a data link. The time for each Figure is shown in milliseconds. The base implementation of rumqtt could not transmit data in the synthetic packet loss scenario.




Figure 5.3: Top contributions to the binary size of the MQicTT client by crate. That is, which libraries contribute the largest footprint.






Figure 5.4: Top contributions to the binary size of the MQuicTT broker by crate. That is, which libraries contribute the largest footprint.



(a) Client binary size reduction.



(b) Broker binary size reduction.

Figure 5.5: The stages of the reduction and the sizes of the QUIC binary at the corresponding stages.

6 | TLS and possible alternatives

As we have shown in Section 5.3, the TLS implementation size remains the main barrier for secure communication using MQTT. Hence, in this section, we analyse what exactly TLS provides and possible lightweight alternatives.

[Description of TLS here]

[Description of KP-ABE here]

[Discussion here]

7 | Conclusion

A | Data

B | MQTT simulation messages

C | Mininet Topology Definition

The full mininet topology definition via the *dump* and *links* command. The Python API definition can also be found on GitHub ¹.

Links:

```
h00-eth1<->h01-eth1 (OK OK) h00-eth2<->h02-eth1 (OK OK)
h00-eth3<->h03-eth1 (OK OK) h00-eth4<->h04-eth1 (OK OK)
h01-eth2<->h02-eth2 (OK OK) h01-eth3<->h03-eth2 (OK OK)
h01-eth4<->h04-eth2 (OK OK) h02-eth3<->h03-eth3 (OK OK)
h02-eth4<->h04-eth3 (OK OK) h03-eth4<->h04-eth4 (OK OK)
h10-eth1<->h11-eth1 (OK OK) h10-eth2<->h12-eth1 (OK OK)
h10-eth3<->h13-eth1 (OK OK) h10-eth4<->h14-eth1 (OK OK)
h11-eth2<->h12-eth2 (OK OK) h11-eth3<->h13-eth2 (OK OK)
h11-eth4<->h14-eth2 (OK OK) h12-eth3<->h13-eth3 (OK OK)
h12-eth4<->h14-eth3 (OK OK) h13-eth4<->h14-eth4 (OK OK)
s0-eth1<->h00-eth0 (OK OK) s0-eth2<->h01-eth0 (OK OK)
s0-eth3<->h02-eth0 (OK OK) s0-eth4<->h03-eth0 (OK OK)
s0-eth5<->h04-eth0 (OK OK) s0-eth6<->s9-eth2 (OK OK)
s1-eth1<->h10-eth0 (OK OK) s1-eth2<->h11-eth0 (OK OK)
s1-eth3<->h12-eth0 (OK OK) s1-eth4<->h13-eth0 (OK OK)
s1-eth5<->h14-eth0 (OK OK) s1-eth6<->s9-eth3 (OK OK)
s9-eth1<->h9-eth0 (OK OK)
```

Listing C.1: Mininet links on the used topology.

Dump:

```
<Host h00:
  h00-eth0:10.0.0.1,h00-eth1:None,h00-eth2:None,h00-eth3:None,h00-eth4:None
  pid=17658>
<Host h01:
  h01-eth0:10.0.0.2,h01-eth1:None,h01-eth2:None,h01-eth3:None,h01-eth4:None
  pid=17660>
<Host h02:
  h02-eth0:10.0.0.3,h02-eth1:None,h02-eth2:None,h02-eth3:None,h02-eth4:None
  pid=17662>
<Host h03:
  h03-eth0:10.0.0.4,h03-eth1:None,h03-eth2:None,h03-eth3:None,h03-eth4:None
  pid=17664>
```

¹Topology - https://github.com/Apolexian/level4-scripts/blob/master/mesh_topo/mesh.py


```

<Host h04:
    h04-eth0:10.0.0.5,h04-eth1:None,h04-eth2:None,h04-eth3:None,h04-eth4:None
    pid=17666>
<Host h9: h9-eth0:10.0.0.6 pid=17668>
<Host h10:
    h10-eth0:10.0.0.7,h10-eth1:None,h10-eth2:None,h10-eth3:None,h10-eth4:None
    pid=17670>
<Host h11:
    h11-eth0:10.0.0.8,h11-eth1:None,h11-eth2:None,h11-eth3:None,h11-eth4:None
    pid=17672>
<Host h12:
    h12-eth0:10.0.0.9,h12-eth1:None,h12-eth2:None,h12-eth3:None,h12-eth4:None
    pid=17674>
<Host h13:
    h13-eth0:10.0.0.10,h13-eth1:None,h13-eth2:None,h13-eth3:None,h13-eth4:None
    pid=17676>
<Host h14:
    h14-eth0:10.0.0.11,h14-eth1:None,h14-eth2:None,h14-eth3:None,h14-eth4:None
    pid=17678>
<OVSSwitch s0:
    lo:127.0.0.1,s0-eth1:None,s0-eth2:None,s0-eth3:None,s0-eth4:None,s0-eth5:None,
    s0-eth6:None pid=17683>
<OVSSwitch s1:
    lo:127.0.0.1,s1-eth1:None,s1-eth2:None,s1-eth3:None,s1-eth4:None,s1-eth5:None,
    s1-eth6:None pid=17686>
<OVSSwitch s9: lo:127.0.0.1,s9-eth1:None,s9-eth2:None,s9-eth3:None pid=17689>
<Controller c0: 127.0.0.1:6653 pid=17651>

```

Listing C.2: Mininet dump on the used topology.

7 | Bibliography

- F. A. Alaba, M. Othman, I. A. T. Hashem, and F. Alotaibi. Internet of Things security: A survey. *J. Netw. Comput. Appl.*, 2017. doi: 10.1016/j.jnca.2017.04.002.
- R. Alena, R. Gilstrap, J. Baldwin, T. Stone, and P. Wilson. Fault tolerance in ZigBee wireless sensor networks. In *2011 Aerospace Conference*, pages 1–15, Mar. 2011. doi: 10.1109/AERO.2011.5747474. ISSN: 1095-323X.
- Bytebeam. rumqtt, Mar. 2020. URL <https://github.com/bytebeamio/rumqtt>. original-date: 2019-10-15T07:49:25Z.
- V. Cerf and R. Kahn. A Protocol for Packet Network Intercommunication. *IEEE Transactions on Communications*, 22(5):637–648, May 1974. ISSN 1558-0857. doi: 10.1109/TCOM.1974.1092259.
- Chromium. QUIC, a multiplexed transport over UDP – The Chromium Projects, 2021. URL <https://www.chromium.org/quic>.
- CoRE. Core working group | ietf core wg, 2016. URL <https://core-wg.github.io/>.
- Eclipse. Eclipse Paho MQTT Rust Client Library, Dec. 2018. URL <https://github.com/eclipse/paho.mqtt.rust>. original-date: 2017-10-03T19:57:53Z.
- L. Eggert. Towards Securing the Internet of Things with QUIC. 2020. doi: 10.14722/diss.2020.23001.
- Ericsson. IoT & Smart manufacturing – Mobility Report, June 2018. URL <https://www.ericsson.com/en/reports-and-papers/mobility-report/articles/realizing-smart-manufact-iot>.
- S. Gupta and N. Lingareddy. Security Threats and Their Mitigations in IoT Devices. 2021. doi: 10.1007/978-3-030-69921-5_42.
- J. Iyengar and M. Thomson. QUIC: A UDP-Based Multiplexed and Secure Transport. Request for Comments RFC 9000, Internet Engineering Task Force, May 2021. URL <https://datatracker.ietf.org/doc/rfc9000>. Num Pages: 151.
- P. Kumar and B. Dezfouli. Implementation and Analysis of QUIC for MQTT. *arXiv:1810.07730 [cs]*, Jan. 2019. URL <http://arxiv.org/abs/1810.07730>. arXiv: 1810.07730.
- B. Lantz and B. Heller. Mininet: An Instant Virtual Network on Your Laptop (or Other PC) – Mininet, 2013. URL <http://mininet.org/>.
- G. Liang, S. R. Weller, J. Zhao, F. Luo, and Z. Y. Dong. The 2015 ukraine blackout: Implications for false data injection attacks. *IEEE Transactions on Power Systems*, 32, 2017. ISSN 08858950. doi: 10.1109/TPWRS.2016.2631891.
- Z. Ling, K. Liu, Y. Xu, C. Gao, Y. Jin, C. Zou, X. Fu, and W. Zhao. IoT Security: An End-to-End View and Case Study. *ArXiv*, 2018.

- N. Matsakis. Baby Steps, Oct. 2019. URL <https://smallcultfollowing.com/babysteps/blog/2019/10/26/async-fn-in-traits-are-hard/>.
- A. Mileva, A. Velinov, L. Hartmann, S. Wendzel, and W. Mazurczyk. Comprehensive analysis of MQTT 5.0 susceptibility to network covert channels. *Computers & Security*, 104: 102207, May 2021. ISSN 0167-4048. doi: 10.1016/j.cose.2021.102207. URL <https://www.sciencedirect.com/science/article/pii/S0167404821000316>.
- M. Miller. MSRC-Security-Research/2019_02 - BlueHatIL - Trends, challenge, and shifts in software vulnerability mitigation.pdf at master · microsoft/MSRC-Security-Research, 2019. URL <https://github.com/microsoft/MSRC-Security-Research>.
- Mozilla. Neqo, an Implementation of QUIC written in Rust, Jan. 2022. URL <https://github.com/mozilla/neqo>. original-date: 2019-02-18T19:20:20Z.
- K. Nepomuceno, I. N. D. Oliveira, R. R. Aschoff, D. Bezerra, M. S. Ito, W. Melo, D. Sadok, and G. Szabo. Quic and tcp: A performance evaluation. volume 2018-June, 2018. doi: 10.1109/ISCC.2018.8538687.
- OASIS. MQTT - The Standard for IoT Messaging, 2014. URL <https://mqtt.org/>.
- Ofcom. UK home broadband performance, measurement period March 2021, Sept. 2021. URL <https://www.ofcom.org.uk/research-and-data/telecoms-research/broadband-research/broadband-speeds/uk-home-broadband-performance-march-2021>.
- J. Postel. User Datagram Protocol. Request for Comments RFC 768, Internet Engineering Task Force, Aug. 1980. URL <https://datatracker.ietf.org/doc/rfc768>.
- S. Previdi, L. Ginsberg, C. Filsfil, A. Bashandy, H. Gredler, and B. Decraene. IS-IS Extensions for Segment Routing. Request for Comments RFC 8667, Internet Engineering Task Force, Dec. 2019. URL <https://datatracker.ietf.org/doc/rfc8667>.
- E. Rescorla. The Transport Layer Security (TLS) Protocol Version 1.3. Request for Comments RFC 8446, Internet Engineering Task Force, Aug. 2018. URL <https://datatracker.ietf.org/doc/rfc8446>. Num Pages: 160.
- J. Roskind. QUIC: Design Document and Specification Rationale - Google Docs, 2012. URL https://docs.google.com/document/d/1RNHkx_VvKWYwg6Lr8SZ-saqsQx7rFV-ev2jRFUoVD34/edit.
- Rust. RAI - Rust By Example, 2021. URL <https://doc.rust-lang.org/rust-by-example/scope/raii.html>.
- SDU. ICTP-SDU: about PingER, Oct. 2013. URL <https://web.archive.org/web/20131010010244/http://sdu.ictp.it/pinger/pinger.html>.
- Z. Shelby, K. Hartke, and C. Bormann. The constrained application protocol (coap). *Rfc* 7252, 2014. ISSN 2070-1721. URL <https://datatracker.ietf.org/doc/html/rfc7252>.
- Statista. Number of IoT devices 2015-2025, 2016. URL <https://www.statista.com/statistics/471264/iot-number-of-connected-devices-worldwide/>.
- S. N. Swamy, D. Jadhav, and N. Kulkarni. Security threats in the application layer in IOT applications. In *2017 International Conference on I-SMAC (IoT in Social, Mobile, Analytics and Cloud) (I-SMAC)*, pages 477-480, Feb. 2017. doi: 10.1109/I-SMAC.2017.8058395.

- T-Mobile. Narrowband iot solution developer protocols guide, 2018. URL https://www.t-mobile.com/content/dam/tfb/pdf/IoT-Solution-Developer-Protocols-Guide.pdf?icid=TFB_TMO_P_19IoT_E3U0LES40BR06VPW19260.
- M. Thomson and S. Turner. Using TLS to Secure QUIC. Request for Comments RFC 9001, Internet Engineering Task Force, May 2021. URL <https://datatracker.ietf.org/doc/rfc9001>. Num Pages: 52.
- I. Vaccari, M. Aiello, and E. Cambiaso. SlowTT: A Slow Denial of Service against IoT Networks. *Information*, 11(9):452, Sept. 2020. ISSN 2078-2489. doi: 10.3390/info11090452. URL <https://www.mdpi.com/2078-2489/11/9/452>.
- M. Weiser. The computer for the 21st century. 1991. doi: 10.1038/SCIENTIFICAMERICAN0991-94.